How Much Fault Tolerance Is in a Fly Brain?

Fault-Injection Analysis of a Whole-Connectome Spiking Controller for Physical Intelligence

Shuchao Gao
Waseda University
Tokyo, Japan

***Abstract*—Whole-brain point-neuron models built from complete synaptic wiring diagrams have begun to serve as controllers in embodied settings, but their behavior under component-level faults has never been measured. We take the released leaky integrate-and-fire (LIF) whole-brain model of *Drosophila melanogaster* (138,639 neurons, 15.1M signed connections, 54.5M synapse contacts) and inject five hardware-motivated fault classes—weight variability, synapse loss, random unit loss, hub loss, and sensory-channel deafferentation—scoring the population response of descending command neurons. Across 36 fault conditions (12 seeded trials each) under sustained escape-grade drive, the connectome absorbs 50% lognormal weight variability, retains 76% of command rate after 50% connection loss, and shows no median change under 10% random unit loss, while loss of the top-1% out-degree hubs produces a 2.5× runaway of global activity rather than silence. An edge-permuted null preserving weight and degree statistics transmits nothing to the command readout, and relative pathway decay is topology-independent: robustness is consistent with threshold amplification on convergence structure, not graph redundancy alone. We distill the results into five candidate design implications for fault-tolerant neuromorphic controllers.**

***Index Terms*—neuromorphic computing, connectome, fault tolerance, fault injection, Drosophila**

I. Introduction

PHYSICAL intelligence—agents whose perception, decision, and action loops close in the physical world under tight energy and latency budgets—demands controllers that remain robust on imperfect hardware [1], [2]. Neuromorphic systems promise the required efficiency [3], [4], and spiking implementations on Loihi 2 and FPGA substrates now cover vision and control workloads [5]–[7]. What remains poorly quantified is how *wiring* determines robustness: engineered spiking networks are designed task-by-task, so their fault tolerance must be re-measured per design [8], [9]. Complete, signed wiring diagrams of whole brains are now public, and point-neuron models built on them reproduce realistic dynamics without training [10], [11]; an evolved connectome is thus a field-tested controller and a reference point for engineered fault tolerance.

The released leaky integrate-and-fire (LIF) whole-brain model of *Drosophila melanogaster* [10][1], which instantiates the complete FlyWire connectome [11], already used targeted computational silencings to identify circuit neurons [10], and we extend that perturbation style into a systematic sweep across *hardware-motivated fault classes*, scored jointly for transmission and activity-envelope compliance. We instrument the visuomotor escape pathway: looming-sensitive visual interneurons (LC4, LPLC2) converge onto descending command neurons (DNp01, DNa02, DNp04, DNb01) whose activation evokes take-off and escape steering [12]. Its failure modes map directly onto hardware fault classes: weight imprecision, broken interconnect, dead cores, dead sensors.

We inject five fault families (Section III-C), each over 12 seeded 600 ms trials, with two controls: an edge-permuted null preserving weights and degree statistics, and an analytic pathway-mass computation. Our contributions: (i) a fault-injection framework mapping hardware failure modes onto a connectome-constrained 138k-neuron LIF model; and (ii) the first tolerance curves of descending-neuron output in a complete connectome (50% weight noise absorbed, 50% connection loss retains 76%, 10% unit loss leaves the median unchanged): top-1% hub loss instead produces a $2.5\times$ *activity runaway* that passes transmission while violating the output envelope, and whole-channel deafferentation is the only catastrophic mode. We further (iii) decompose why graph redundancy alone explains neither specificity nor graceful degradation (threshold amplification as a candidate account) and (iv) distill five candidate design implications for fault-tolerant controllers.

II. Related Work

The release of complete wiring diagrams for the adult *Drosophila* brain [11] has shifted connectome research from static graph analysis to executable models spanning visual system convolutions [13], neuron-level response prediction [14], learned tasks and whole-body locomotion on whole-CNS wiring [15]–[17], and biofidelic processing units [18], with cell-type studies resolving descending-neuron input sharing [19], [20]. Closest to our substrate, the released LIF whole-brain model (footnote 1) reproduces global activity and has been ported to Loihi 2 [5]. None of these works varies the integrity of the controller; we treat the connectome controller as a fixed plant and sweep internal fault modes.

Graph-theoretic lesion studies score the structural attack tolerance of brain networks [21], not dynamics; spiking-network weight-noise robustness is mostly characterized on small trained architectures [8], [9], with adaptive mitigation now studied at the edge [22]; purpose-built fault-injection frameworks for trained SNNs—SpikeFI [23], SpikingJET [24]—automate this but score task accuracy on human-designed

[1]https://github.com/philshiu/Drosophila_brain_model

networks; follow-up work covers STDP-trained robustness and population-coded redundancy [25], [26]. Hardware platforms expose complementary constraints: on-chip variability co-design [6], FPGA precision–capacity trade-offs [7], permanent interconnect faults [27], edge energy–robustness envelopes [3], [28], and trained-CNN dependability [29]. Benchmarking efforts omit fault-injection axes [30], and formalizations of neuromorphic control do not yet treat degradation [31]. Insect-inspired collision-aversion circuits are robust by construction, but small and hand-designed [32], [33]. What a complete connectome uniquely enables is tolerance measurement at the scale of an entire evolved controller; these measurements suggest hypotheses for subsequent hardware validation.

III. Fault-Injection Framework

A. Whole-Brain Spiking Substrate

We adopt the released whole-brain LIF model of *Drosophila* [10] (Section I), implemented in Brian2 [34] and built on the FlyWire connectome [11]. Every neuron is a current-based LIF unit (threshold -45 mV, reset -52 mV, membrane time constant 20 ms, refractory 2.2 ms, axonal delay 1.8 ms). Each connection in the released table becomes a delta-conductance synapse whose jump is 0.275 mV per counted contact and decays with a 5 ms synaptic time constant, signed by the presynaptic cell's neurotransmitter; the table yields 15,091,983 connections between 138,639 neurons (54.5M contacts), 40% inhibitory. As in the released model, baseline firing is zero: the only external input is the stimulus Poisson group over the driven set, so every recorded spike traces back to the stimulated cells. Our engine restores the full network state at the start of every trial, so all conditions share identical initial states, while the stimulus draws a fresh realization each trial: reported dispersion includes stimulus variability. One 600 ms trial runs in $\approx$3.5 s on one CPU core.

B. Task: Visuomotor Escape

The instrumented pathway (Fig. 1) is the fly's looming-triggered escape circuit: looming-sensitive visual projection neurons LC4 (104) and LPLC2 (210) converge on descending neurons [12], [32]. Because collision-sensing models treat looming as sustained high-rate drive, we drive the sensory cells—and, as a control, an equal-sized random set—with independent 150 Hz Poisson trains gated per trial, rather than modeling looming kinetics. The primary readout is the population rate of the command set {DNp01, DNa02, DNp04, DNb01} (8 neurons); all 1,299 descending neurons serve as a secondary readout. Calibration drove the command population to 94.6 Hz (interdecile 93.6–95.8 Hz, first-spike latency 4.4 ms; 10-trial set) under escape stimulation but only 8.2 Hz (6.4–9.2 Hz) under random stimulation; $\theta = 51.4$ Hz (the 94.6/8.2 Hz midpoint); every pass/fail classification below is unchanged for thresholds between the observed trial extremes (9.2–63.1 Hz); Section IV uses the 12 sweep trials.

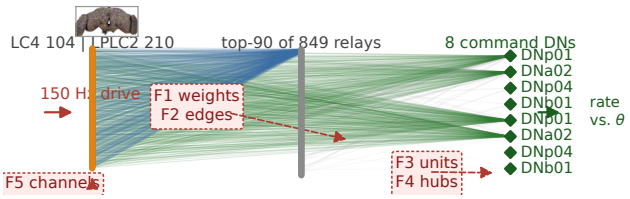


Fig. 1. The instrumented escape pathway, drawn from the connectome itself. Left: the 314 driven sensory cells (LC4, LPLC2); center: the 90 top-ranked of 849 two-hop relays; right: the 8 command DNs. Line width and opacity scale with synapse count (594 direct sensory→DN edges); dashed red marks the fault-injection points (F5 channels; F1/F2 synapses; F3/F4 units). Inset: reconstructed adult brain, adapted from Dorkenwald et al. [35] (CC BY 4.0).

C. Fault Taxonomy

Five fault families cover the hardware-relevant failure modes. F1 multiplies every synaptic weight by an i.i.d. mean-preserving lognormal factor ($\sigma \in \{0.05, 0.1, 0.2, 0.5\}$), modeling analog weight variability and bit-level imprecision. F2 deletes a random fraction $f \in \{0.05, 0.10, 0.25, 0.50\}$ of connections ($\approx f$ of contacts in expectation), modeling interconnect or memristor failures. F3 silences all output synapses of a random fraction $f \in \{0.01, 0.02, 0.05, 0.10\}$ of units, excluding input and readout cells (unit removal in this delta-synapse model), modeling dead cores. F4 applies the same operation to the top 1%, 2%, or 5% of units ranked by total outgoing contact count (again excluding input and readout), modeling failure of high-fanout components; these cells carry 9–22% of all connections (1%, 2% and 5% sets, respectively), a targeted, correlated version of F3. F5 zeroes all outgoing synapses of one (LC4 or LPLC2) or both sensory channels, modeling sensor or afferent-path failure. All faults are applied after state restoration; F4 and F5 act on fixed cell sets, while F1–F3 redraw their realization each trial from SHA-256 seeds over (phase, condition, trial)—12 independent draws per condition, exactly reproducible from the archived seeds (the stimulus stream is not re-seeded, so re-runs are statistically equivalent rather than bit-identical).[2]

D. Controls

Two controls isolate what the biological wiring contributes. The first is an edge-permuted null: every synapse keeps its presynaptic partner, weight, and sign, but its postsynaptic target is randomly reassigned. This preserves out-degrees, weights, and the global in-degree histogram while destroying which inputs converge onto each target—and with them all multi-hop structure; the identical sweep runs on it. The second is analytic: for both graphs we compute the *pathway mass* $m_k = \mathbf{1}_{\mathrm{DN}}^{\top} A^k \mathbf{1}_{\mathrm{in}}$, the synapse-weighted total of all k-hop directed paths from LC4∪LPLC2 into the full descending layer (A holds contact counts, i.e., absolute signed strengths; m_k counts k-step walks of either sign—a wiring-density proxy, not signal flow), on the intact graph and after the F2 deletion

[2]Code, connectome-derived edge lists, and per-trial records: https://github.com/Scophield/drosophila-fault-injection.

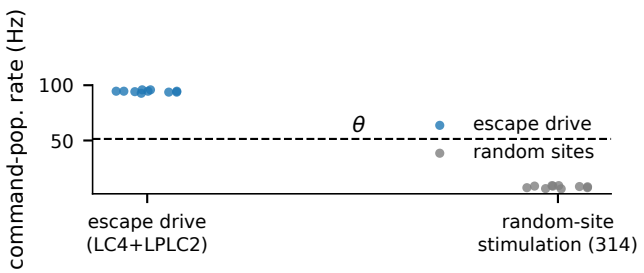


Fig. 2. Readout specificity. Command rate under escape stimulation (LC4+LPLC2) versus matched random-site stimulation (314 cells); dashed line: θ.

levels; comparing its decay with the output decay separates topological redundancy from dynamical amplification.

IV. Results

A. Baseline: The Connectome Is a Specific, Low-Latency Controller

Driving LC4 and LPLC2 at 150 Hz produces a command rate of 94.7 Hz (median over 12 trials; interdecile 94.0–95.6 Hz) with 4.05 ms median first-spike latency, far above θ (Fig. 2). An equal-sized random set, driven identically, yields only 8.2 Hz on the same readout while producing *more* total activity (293k vs. 40k spikes): the readout is pathway-specific, not generic arousal. When both sensory channels are silenced ($F5_{both}$), the command rate collapses to zero in every trial; the residual 21.2k spikes are, by construction, the stimulated sensory cells themselves—the rest of the brain, initialized at zero baseline, receives no input—so the score measures transmission through the instrumented pathway, not global silencing.

B. Tolerance Curves: Noise Is Absorbed, Disconnection Degrades Gracefully

Figure 3 shows dose–response curves for the two most consequential families; F1 and F3 responses are flat in median (Table I). Weight variability (F1) leaves the median unchanged: at σ=0.5 the rate is 94.2 Hz (healthy: 94.7 Hz), the trial dispersion widens (interdecile 90.4–100.8 vs. 94.0–95.6 Hz), and total activity rises only from 39.7k to 41.6k spikes. Random connection deletion (F2) degrades gracefully: 5% and 10% loss leave the rate at 91.9 and 93.1 Hz, 25% loss retains 84.5 Hz (89.2% of baseline), and even 50% loss retains 72.4 Hz (76.5%), with all 12 trials passing transmission in every condition (weakest trials 63.1 and 68.1 Hz—above θ by $\geq$22%). Disconnection's cost appears predominantly in latency: 4.05 to 5.4 ms at 50% loss (+33%), consistent with recruitment of longer residual multi-hop routes; which dominates a task-level error budget is workload-dependent. Random unit loss (F3) up to 10%—nearly 14,000 neurons with all their synapses—leaves the median unchanged (95.1 vs. 94.7 Hz): command drive integrates over hundreds of converging populations, and no tested random slice is load-bearing. The wide 1,299-neuron descending layer tracks the same pattern (2,941 spikes healthy vs. 1,918 at 50% loss), so the graded response is a property of command encoding, not of the eight-neuron readout. Aggregation also masks composition: four of the eight command cells carry 96.6% of healthy readout spikes, and at 50% connection loss the smallest contributors collapse (158→10, 26→0) while silent cells are recruited (1→58, 0→21); the curves are thus a descending-population response, not behavior.

C. Hub Loss Produces Runaway, Not Silence

The dose–response of hub removal (F4) is non-monotonic and qualitatively different from every distributed fault (Figs. 3b, 4b). Removing the top 1% of cells by outgoing synapse count (1,386 neurons, 9.1% of connections) *increases* the command rate to 122.3 Hz (+29%, interdecile 121.1–123.9 Hz across trials) and more than doubles global activity (99.9k spikes, 2.52× healthy; Mann–Whitney vs. healthy $p < 10^{-4}$, as for $F2_{50}$ and $F5_{both}$); at 2% the rate stays elevated (117.1 Hz, activity 42.4k); at 5% (22.4% of connections) both metrics return to near-baseline (94.2 Hz, 36.5k). A matched-scale contrast rules out a size effect: random 10% unit ablation disconnects 10.0% of connections (more than the hub set's 9.1%) yet leaves the readout unchanged (95.1 Hz). Hub lesions do not starve the pathway of drive—they remove broadly connected inhibitory regulation, and spared circuitry over-drives the command layer: 62% of removed hub output connections are inhibitory (40% network-wide). We therefore score *transmission* (rate $\geq \theta$) and an *activity envelope* (rate and total activity within 2× of healthy) separately: hub loss at 1% passes transmission while violating the envelope in all 12 trials (persisting for any multiplier between 1.5× and 2×); every distributed fault passes both; $F5_{both}$ fails both. The envelope violation is the engineered system's dangerous mode: output exceeds its envelope while every per-unit monitor reads "active." The same lesion in the permuted null produces no such signature, tying the runaway to biological hub structure, not degree heterogeneity.

D. What the Wiring Contributes: Null Network and Pathway Analysis

The permuted null dissociates specificity from redundancy: out-degrees, weights, and the in-degree histogram survive, convergence is reassigned, and it transmits *nothing* to the command layer (median 0.0 Hz in every condition; no null trial reaches θ); its broad 1,299-neuron readout falls to background (45.5 spikes per trial vs. 2,941 healthy) and global activity to about 60% of the biological level (23.8k vs. 39.7k)—little more than the 21.2k the driven cells produce with all outputs severed ($F5_{both}$): the permuted network receives the same drive, but only ∼2.6k spikes propagate and none converge onto descending output. Pathway mass quantifies the convergence advantage: 5.8× (1-hop) and 17.7× (4-hop) more route mass into the descending layer than the null. Yet *relative* pathway-mass decay is nearly identical in both graphs (at f=0.5, k=8: 0.3% vs. 0.4% of baseline; Fig. 4a), while output degrades far more slowly than pathway mass (76% at

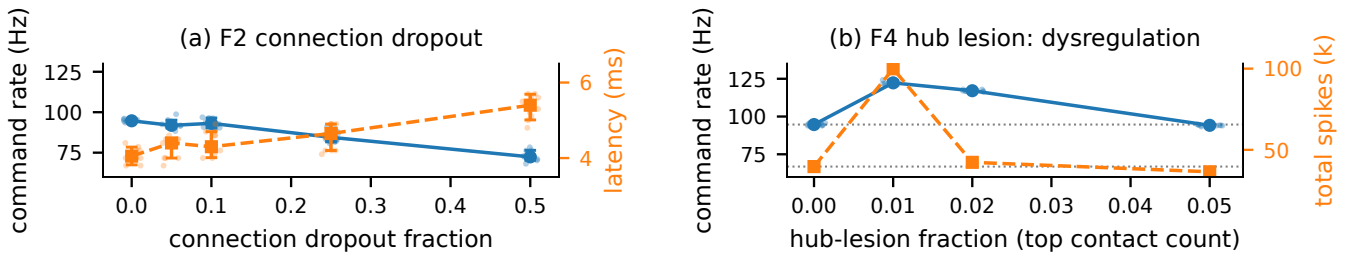


Fig. 3. Tolerance curves of descending-neuron output (medians over 12 trials; blue: rate; orange: latency (a) / total spikes (b); dotted grey: healthy medians). (a) Connection dropout: 76% of command rate retained at 50% loss, latency +33%. (b) Hub lesion: dysregulation.

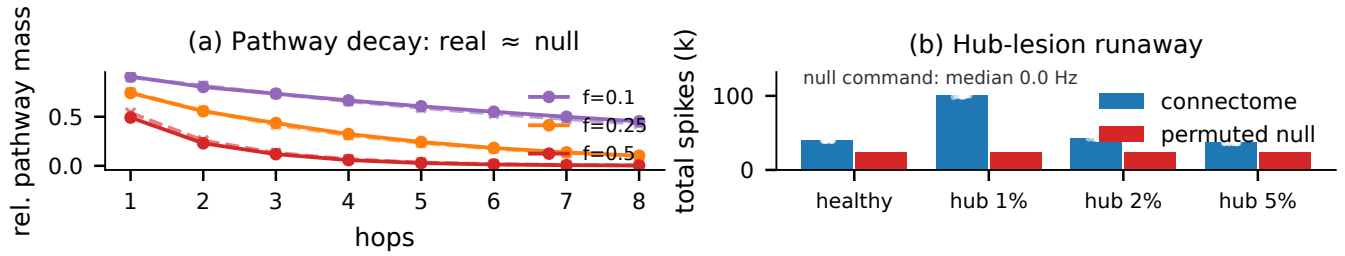


Fig. 4. Topology contrast. (a) Route mass under edge loss, normalized per graph (connectome solid, null dashed): relative decay nearly identical, absolute mass 5.8–17.7× higher. (b) Activity under hub lesions: runaway only in biological wiring; null readout 0.0 Hz (median).

f=0.5 against 5.8% four-hop route mass)—ruling out graph redundancy alone as the mechanism. A candidate account is threshold amplification: with N converging afferents and per-synapse survival p, aggregate drive scales as Np while k-hop route mass decays as p^k in expectation (6.25% at k=4, p=0.5), and the saturating threshold compresses drive changes into small rate changes. Retention above proportional scaling (76.5% at p=0.5, 89.2% at p=0.75) is consistent with this account but does not establish it: in an excitatory–inhibitory network even linear signal flow need not scale with p, since deleting inhibitory connections offsets excitatory loss. Disentangling disinhibition, direct sensory→DN edges, and threshold amplification requires E/I-balanced deletion and drive-strength sweeps.

V. Candidate Design Implications for Embodied Neuromorphic Controllers

Table I distills five candidate design implications. P1 and P2 follow from the flat F1 and F3 medians: computation lives in distributed convergence, not precise magnitudes, consistent with hardware observations [6], [8]. P3: disconnection costs are mostly latency, so real-time loops should budget for the median latency growth that Fig. 3a reports. P4: removing the highest-fanout units is the only *amplifying* failure mode, so cap high-degree influence and monitor output-envelope violation rather than liveness. P5: each sensory channel alone suffices (68–79%); front-end duplication buys more margin than internal redundancy.

One methodological lesson: identical relative pathway decay, yet only the biological graph transmitting a decodable command, argues against certifying embodied controllers by topology alone; fault-injection campaigns belong in pre-deployment stages [30], [31].

TABLE I
Candidate design implications from connectome fault injection, with quantitative anchors.

#	Implication	Anchor (this study)
P1	Buy robustness with redundancy of wiring, not precision of weights	σ=0.5 lognormal noise: 99.4% retention
P2	Distribute command integration; no tested slice load-bearing	10% unit loss: 100.4% retention
P3	Budget latency for disconnection faults	50% connection loss: 76% rate, +33% latency
P4	Regulate high-fanout units; monitor envelope, not liveness	top-1% hub loss: 2.5× activity runaway
P5	Duplicate the sensory interface before the network	single-channel loss 68–79% median; both channels 0%

VI. Limitations and Conclusion

Our measurements inherit the LIF abstraction—single-compartment units, no plasticity, Poisson drive, one drive rate; doses kept the background regime intact, the hub analysis is descriptive, and correlated damage and recovery are excluded. Under the tested drive (sustained, well-above-threshold), the controller keeps its median command rate under weight noise, unit death, and 50% connection loss; its dangerous failure is hub-loss dysregulation.

REFERENCES

[1] Y. Yuan, Y. Huang, X. Yao, Y. Li, S. Zhang, L. Han *et al.*, "Embodied-r1.5: Evolving physical intelligence via embodied foundation models," arXiv:2606.11324, 2026.

[2] W. Guo, H. Zhang, P. Li, T. Cai, Z. Chen, Y. Guo *et al.*, "Neurovla: A brain-inspired embodied intelligence for fluid and fast reflexive robotics control," arXiv:2601.14628, 2026.

[3] K. Katti, P. Chaudhari, and D. Jariwala, "Neuromorphic computing for low-power artificial intelligence," arXiv:2604.04727, 2026.

[4] M. van Gerven, "Neuromorphic intelligence," arXiv:2509.11940, 2025.

[5] F. Wang, B. H. Theilman, F. Rothganger, W. Severa, C. M. Vineyard, and J. B. Aimone, "Neuromorphic simulation of drosophila melanogaster brain connectome on loihi 2," arXiv:2508.16792, 2025.

[6] S. Arfa, B. Vogginger, and C. Mayr, "Hardware-aware fine-tuning of spiking q-networks on the spinnaker2 neuromorphic platform," ACM ICONS 2025 - International Conference on Neuromorphic Systems, 2025.

[7] A. Gautam, P. Date, S. Kulkarni, R. Patton, and T. Potok, "Neurocorex: An open-source fpga-based spiking neural network emulator with on-chip learning," arXiv:2506.14138, 2025.

[8] T. K. T. Thieu and R. Melnik, "Effects of noise on leaky integrate-and-fire neuron models for neuromorphic computing applications," arXiv:2202.09482, 2022.

[9] A. Yunusoglu, D. Le, M. Isik, I. C. Dikmen, and T. Karadag, "Neuromorphic circuits with spiking astrocytes for increased energy efficiency, fault tolerance, and memory capacitance," arXiv:2502.20492, 2025.

[10] P. K. Shiu, G. R. Sterne, N. Spiller, R. Franconville, A. M. Sandoval, J. Zhou *et al.*, "A drosophila computational brain model reveals sensorimotor processing," *Nature*, vol. 634, no. 8032, pp. 210–219, 2024.

[11] A. Lin, R. Yang, S. Dorkenwald, A. Matsliah, A. R. Sterling, P. Schlegel *et al.*, "Network statistics of the whole-brain connectome of drosophila." in *Nature*, 2024.

[12] J. H. Simpson, "Descending control of motor sequences in drosophila," *Current Opinion in Neurobiology*, vol. 84, p. 102822, 2023.

[13] J.-N. Liew, S. Lin, B. Chen, W. Zhang, X. Zhu, W. Zhang *et al.*, "Connectome-based modelling reveals orientation maps in the drosophila optic lobe," arXiv:2609.01330, 2026.

[14] J. Xie, R. Ren, X. Zhou, A. Zheng, J. Zhu, W. Jiang *et al.*, "Visual function profiles via multi-path aggregation reveal neuron-level responses in the drosophila brain," arXiv:2512.06934, 2025.

[15] J. Cho, M. Kang, and D. Kim, "A drosophila whole-connectome network can learn human-designed cognitive tasks," arXiv:2610.10014, 2026.

[16] S. Senge, "A connectome test of the fly hashing algorithm," arXiv:2610.09114, 2026.

[17] Z. Jin, Y. Zhu, C. Zhang, and Y. Sui, "Whole-brain connectomic graph model enables whole-body locomotion control in fruit fly," arXiv:2602.17997, 2026.

[18] S. Yu, Z. Qin, T. Liu, B. Xu, R. J. Vogelstein, J. Brown *et al.*, "Biological processing units: Leveraging an insect connectome to pioneer biofidelic neural architectures," arXiv:2507.10951, 2025.

[19] X. Zhang, "Synaptic placement reflects shared input in drosophila descending neurons," arXiv:2610.00690, 2026.

[20] N. Drummond, A. Zhao, and A. Borst, "Population morphology implies a common developmental blueprint for drosophila motion detectors." in *PLoS computational biology*, 2026.

[21] H. Aerts, W. Fias, K. Caeyenberghs, and D. Marinazzo, "Brain networks under attack: robustness properties and the impact of lesions," *Brain*, vol. 139, no. 12, pp. 3063–3083, 2016.

[22] S. Luca, T. P. Xiao, F. S. Chance, S. Agarwal, C. Teeter, and G. W. Chapman, "Learned adaptive properties for mitigation of weight perturbations in embedded spiking networks," *Frontiers in Neuroscience*, vol. 20, p. 1766765, 2026.

[23] T. Spyrou, S. Hamdioui, and H.-G. Stratigopoulos, "Spikefi: A fault injection framework for spiking neural networks," arXiv:2412.06795, 2024.

[24] A. B. Göğebakan, E. Magliano, A. Carpegna, A. Ruospo, A. Savino, and S. D. Carlo, "Spikingjet: Enhancing fault injection for fully and convolutional spiking neural networks," in *2024 IEEE 30th International Symposium on On-Line Testing and Robust System Design (IOLTS)*. IEEE, 2024, pp. 1–7.

[25] Z. Jouni and H.-G. Stratigopoulos, "STDP-trained spiking neural network reliability assessment through fault injections," in *2025 IEEE 31st International Symposium on On-Line Testing and Robust System Design (IOLTS)*. IEEE, 2025, pp. 1–8.

[26] A. Gleyo and B. Girau, "Population coding to improve fault tolerance of neuromorphic networks in regression tasks," in *2026 Design, Automation and Test in Europe Conference (DATE)*. IEEE, 2026, pp. 1–7.

[27] G. Chessa, Y. Yang, E. Jones, M. Favalli, and D. Bertozzi, "Cross-layer analysis of permanent faults in asynchronous NoCs for neuromorphic computing," in *Proceedings of the International Conference on Neuromorphic Systems*. ACM, 2026, pp. 31–38.

[28] O. Y. L. Imanov, D. U. Kulali, T. Yilmaz, D. Erisken, and R. I. Turhan, "Energy-efficient neuromorphic computing for edge ai: A framework with adaptive spiking neural networks and hardware-aware optimization," arXiv:2602.02439, 2026.

[29] J. Gracia-Morán, J. C. Baraza-Calvo, J.-C. Ruiz, D. de Andrés, D. Gil-Tomás, and L.-J. Saiz-Adalid, "Dependability analysis of a spiking convolutional neural network (S-CNN)," *Lecture Notes in Computer Science*, pp. 147–161, 2026.

[30] J. Yik, K. V. den Berghe, D. den Blanken, Y. Bouhadjar, M. Fabre, P. Hueber *et al.*, "Neurobench: A framework for benchmarking neuromorphic computing algorithms and systems," arXiv:2304.04640, 2023.

[31] T. Medvedeva, A. Franci, and F. Castaños, "Formalizing neuromorphic control systems: A general proposal and a rhythmic case study," arXiv:2506.10203, 2025.

[32] F. Shuang, Y. Zhu, Y. Xie, L. Zhao, Q. Xie, J. Zhao *et al.*, "Opplod: the opponency based looming detector, model extension of looming sensitivity from lgmd to lplc2," arXiv:2302.10284, 2023.

[33] Z. Qin, Q. Fu, J. Peng, and S. Yue, "Dynamic neural field modeling of visual contrast for perceiving incoherent looming," arXiv:2504.04551, 2025.

[34] D. F. M. Goodman, "The brian simulator," *Frontiers in Neuroscience*, vol. 3, no. 2, pp. 192–197, 2009.

[35] S. Dorkenwald, A. Matsliah, A. R. Sterling, P. Schlegel, S.-C. Yu, C. E. McKellar *et al.*, "Neuronal wiring diagram of an adult brain." *Nature*, vol. 634, no. 8032, pp. 124–138, 2024.